

A Low-Cost Single-Probe Differential-pH Aeration Monitor for Seawater Carbonate Alkalinity: Systematic-Error Decomposition and Sub-0.25 dKH Accuracy via Sealed Common-Headspace Equilibration

nowforyou*¹

¹REEF FAMILY

June 22, 2026

Abstract

Carbonate alkalinity (“KH”, expressed in degrees, dKH) is the single most control-critical water-chemistry parameter in reef aquaria, yet affordable continuous monitors are scarce and reagent titration is laborious. We describe an open, reagent-free monitor that infers alkalinity from the equilibrium pH a seawater sample reaches after aeration to atmospheric CO₂, using a *single* pH probe to measure an aquarium sample (“tank”) and a known-alkalinity reference (“ref”) in immediate succession. Because dKH is computed from the pH *difference* ($\text{KH}_{\text{tank}} = \text{KH}_{\text{ref}} \cdot 10^{\Delta\text{pH}}$), absolute probe drift and room-CO₂ excursions are common-mode signals that cancel. Over an nine-day instrumented campaign (28 logged measurements across five firmware/method revisions) we observed a persistent uncorrected error of -0.92 dKH and, rather than zero-point calibrating it away, decomposed it. Three interventions—(i) plateau-following measurement to true gas equilibrium, (ii) a single *sealed common headspace* shared by sample, reference and pump intake, and (iii) an actively aerated 5 L reference anchor that pins the headspace p_{CO_2} —collapsed the dawn-CO₂ error and the diurnal evaporation confound, and shrank measurement scatter to $\text{SD} = \pm 0.08$ dKH (about one-third of the ± 0.25 dKH design goal). A sequence of null experiments (identical water in both vessels, true $\Delta\text{pH} = 0$) then localized the residual: a *complete null* returned $\Delta\text{pH} = -0.003$, $\text{dKH} = 8.387$, i.e. an instrument offset of only -0.06 dKH. The remaining -0.86 dKH was traced to alkalinity loss (CO₂ driven CaCO₃ precipitation / biofilm) in the days-old, small, open tank aliquot—a sample-handling artifact, not a measurement fault. Re-running the null on the same aged aliquot a day later watched it decay from -0.06 to -0.22 dKH with the electronics untouched, confirming the attribution in the time domain. The instrument therefore meets the ± 0.25 dKH goal uncorrected, with zero-point calibration deemed unnecessary. Our independently obtained error structure (room-CO₂ sensitivity, bubble-quality dependence, sealing as the fix) reproduces the published behaviour of a commercial counterpart, providing external validation of both the diagnosis and the remedy.

Keywords: reef aquarium; carbonate alkalinity; total alkalinity; CO₂ equilibration; differential pH; low-cost instrumentation; common-mode rejection; CaCO₃ precipitation.

*Correspondence: taeseokyi@gmail.com

1 Introduction

Stony corals and other calcifying reef organisms deposit calcium carbonate at a rate that continuously consumes carbonate alkalinity. In a closed reef aquarium the resulting alkalinity drawdown is fast (often $0.5\text{--}1.5\text{ dKH day}^{-1}$) and swings of more than $\sim 1\text{ dKH}$ are implicated in coral tissue stress. Alkalinity (A_T , conventionally reported by hobbyists as “KH” in German degrees, $1\text{ dKH} \approx 0.357\text{ meq L}^{-1}$) is therefore the parameter most worth measuring *often*.

The accessible options are unsatisfying for continuous use. Manual reagent titration (e.g. commercial drop kits) is accurate but labour-intensive and discrete. Colorimetric robots automate titration but consume reagent and are expensive. Direct potentiometric pH alone is *not* an alkalinity measurement, because seawater pH depends jointly on A_T and p_{CO_2} , and aquarium p_{CO_2} varies strongly over the day (room ventilation, photosynthesis, respiration).

A reagent-free alternative exploits the fact that if a sample is aerated until its p_{CO_2} equilibrates with the surrounding air, the equilibrium pH becomes a single-valued function of alkalinity at fixed temperature and salinity. A commercial product (AquaWiz) is built on this principle and is used by reef hobbyists, but community reports document recurring inaccuracy tied to room- CO_2 fluctuation and to aeration/bubble quality. The contribution of this paper is not the principle but a *quantitative diagnosis*: starting from a working DIY build, we instrument a nine-day campaign, expose a -0.92 dKH raw error, and show by construction (null experiments) that essentially all of it is a sample-handling artifact rather than an instrument limitation—reducing the true instrument offset to -0.06 dKH and meeting a $\pm 0.25\text{ dKH}$ goal without calibration. Along the way we establish a small set of design rules (sealed common headspace; plateau-following to true equilibrium; an actively aerated reference anchor; a two-gate plateau criterion) that we expect to generalize to any single-probe differential-pH alkalinity monitor.

2 Measurement Principle

2.1 Aeration makes equilibrium pH a proxy for alkalinity

At fixed temperature T and salinity S , the seawater carbonate system has two degrees of freedom; fixing any two of $\{A_T, \text{DIC}, p_{\text{CO}_2}, \text{pH}\}$ determines the rest. Aeration with a large excess of room air drives the dissolved CO_2 toward Henry’s-law equilibrium with the gas phase, $[\text{CO}_2]_{\text{aq}} = K_H p_{\text{CO}_2}$, so that $p_{\text{CO}_2}(\text{water}) \rightarrow p_{\text{CO}_2}(\text{air})$. Once p_{CO_2} is clamped by the atmosphere, the equilibrium pH is a single-valued function of alkalinity:

$$\text{pH}_{\text{eq}} = f(A_T, p_{\text{CO}_2}, T, S). \quad (1)$$

With p_{CO_2} , T , S held common, the *only* variable that can move pH_{eq} is A_T . Measured equilibrium pH is therefore a proxy for alkalinity.

2.2 Why the differential cancels room CO_2

The instrument never relies on the absolute value of p_{CO_2} . It aerates the reference and the tank sample to the *same* atmospheric ceiling (in the final design, by aerating both against one shared headspace simultaneously, Sec. 4) and reports the *difference* of their equilibrium pH. Because the absolute p_{CO_2} term is common to both samples, it cancels in the difference: a room- CO_2 excursion is a common-mode signal that drops out. What survives the subtraction is the true alkalinity difference plus electrode noise ($\pm 0.005\text{--}0.01\text{ pH}$). This is the central robustness argument of the method, and the reason a *single* probe (rather than two matched probes) is used: a single sensor’s

slow asymmetry-potential and reference-junction drift load identically onto both readings and cancel in ΔpH .

2.3 Over-aeration is impossible; only *insufficient* aeration errs

Equation (1) has a hard ceiling at $p_{\text{CO}_2}(\text{air})$: once the water reaches atmospheric equilibrium, further aeration cannot raise pH. There is consequently no “over-aeration” failure mode. Error enters only through *insufficiency*: if residual CO_2 remains in one vessel, that side’s pH is unfairly depressed and a spurious ΔpH appears. The differential’s assumption is precisely that ref and tank share the same p_{CO_2} ; broken symmetry breaks the cancellation. This motivates measuring to a genuine plateau rather than for a fixed time (Sec. 4).

2.4 Volume independence at full equilibrium

Equation (1) contains no volume term (Henry’s law is intensive). Hence even a large volume mismatch—in our final null, 100 mL of tank water against a 5 L reference—produces *zero* differential error provided both vessels reach *full* equilibrium. The familiar “equal volumes for symmetry” rule applies only under *partial* equilibration, where unequal volumes sit at different points on the first-order approach curve $C(t) - C^* = (C_0 - C^*) e^{-t/\tau}$. The true requirements are just (a) both fully equilibrated and (b) matched p_{CO_2} and T . The larger volume is rate-limiting, with a time constant

$$\tau \propto \frac{V}{K_H Q_{\text{gas}}} \beta \sim \frac{1}{\text{VVM}} \beta, \quad (2)$$

where $\text{VVM} = Q_{\text{gas}}/V$ is gas volume per liquid volume per minute and β is a seawater chemical factor (only trace $[\text{CO}_2]_{\text{aq}}$ is strippable; the Revelle buffering raises τ by $\sim 8\text{--}15\times$). Equation (2) is the basis for the “increase aeration to shorten the measurement” lever discussed later, and for the choice to miniaturize the rate-limiting reference.

2.5 The differential dKH formula

The firmware (Arduino Nano / ATmega328P; pH via DFRobot Gravity probe + ADS1115 16-bit ADC) digitizes one probe, reading reference then tank pH, and computes

$$\boxed{\text{KH}_{\text{tank}} = \text{KH}_{\text{ref}} \cdot 10^{-(\text{pH}_{\text{ref}} - \text{pH}_{\text{tank}})} = \text{KH}_{\text{ref}} \cdot 10^{\Delta\text{pH}}} \quad (3)$$

with $\Delta\text{pH} \equiv \text{pH}_{\text{tank}} - \text{pH}_{\text{ref}}$ and $\text{KH}_{\text{ref}} = 8.448 \text{ dKH}$ an EEPROM-stored *anchor* constant (set by a `calref` command after evaporation top-off; *not* a per-run measurement—this is why the reference column reads a constant 8.448 in every log line). Equation (3) depends only on ΔpH ; absolute pH cancels exactly. Differentiating,

$$\left. \frac{d\text{KH}_{\text{tank}}}{d\Delta\text{pH}} \right|_{\Delta\text{pH}=0} = \text{KH}_{\text{ref}} \ln 10 \approx 19.45 \text{ dKH pH}^{-1}, \quad (4)$$

i.e. 0.01 pH of uncanceled difference maps to $\approx 0.19 \text{ dKH}$. This high sensitivity is what makes both the cancellation and the plateau criterion matter. Temperature compensation is applied per reading by a Nernst relation $\text{pH}_{\text{corr}} = 7.0 + (\text{pH}_{\text{raw}} - 7.0) (273.15 + T_{\text{cal}}) / (273.15 + T_{\text{meas}})$, nominally zero-error over $15\text{--}35^\circ\text{C}$. Each pH reading is a trimmed mean of 64 samples (discard the top and bottom 16, average the middle 32) to reject aeration-bubble spikes.




Figure 1: Schematic of the sealed closed-loop monitor. A single pH probe in the measurement chamber reads the tank aliquot and the reference in turn; four bidirectional peristaltic pumps (M1–M4) move liquid between the main aquarium, holding beaker, measurement chamber, reference reservoir, and KCl store; two air pumps (**ron**) aerate the reference and the chamber simultaneously. The reference, chambers and pump intakes share one sealed headspace (≈ 4 L), so sample and reference see identical p_{CO_2} and evaporation is suppressed; the controller sits outside. Between runs the probe rests in 3 M KCl (M3).

3 Apparatus

The build (Table 1, Fig. 1) is a closed fluidic loop around four bidirectional peristaltic pumps, two USB air pumps, a single pH probe, and a 1-Wire temperature sensor, all driven by an Arduino Nano over a 9600-baud Bluetooth (HC-06) serial link to a host PC.

Fluidic topology (current). All pumps are bidirectional (**f** = forward, **b** = reverse):

- **M1:** main aquarium $\leftrightarrow$ *holding* beaker;
- **M2:** holding $\leftrightarrow$ *measurement* chamber;
- **M3:** 3 M KCl beaker $\leftrightarrow$ measurement chamber (probe soak / storage);
- **M4:** 5 L reference $\leftrightarrow$ measurement chamber.

Routing the tank sample through an intermediate holding beaker lets it be temporarily *parked* while the reference is measured, minimizing the between-phase delay (Sec. 4). Run times are encoded as the integer in **mXf:NN** / **mXb:NN** (seconds); reverse runs ~ 8 – 10 s longer than forward to fully clear the hose (current values **m1f:70/m1b:82**, **m2f:60/m2b:68**, **m3f:60/m3b:68**, **m4f:60/m4b:70**; the 5 L hose is longest, hence **m4b:70**).

Control commands. Air and motor enables are independent lines: **ron** turns *both* air pumps on (used only while liquid is stationary—bubble cleaning and measurement); **ton** enables the PWM motor controller (*required* before any pump will turn); **airoff** drops both. A load-bearing invariant follows: because **airoff** also disables the PWM controller, every liquid move must be preceded by **airoff** $\rightarrow$ **ton**; the last air command before a motor must be **ton**. Measurement primitives are **tank** / **ref** (sample and store one pH reading) and **calkh** (apply Eq. (3)). (Firmware retains legacy solenoid-era pin names from an earlier hardware revision; pin behaviour is as described.)

Sealed common headspace (required). The 5 L reference, the measurement chamber, and the pump intake share *one* sealed headspace. This is not an optional refinement: it (i) guarantees tank

Table 1: Bill of materials (current build).

Subsystem	Part	Qty
Microcontroller	Arduino Nano V3.0 (ATmega328P)	1
pH probe	DFRobot Gravity pH V2 (SEN0161-V2), single glass electrode	1
ADC	DFRobot Gravity I ² C ADS1115 (16-bit)	1
Temperature	DS18B20 waterproof (1-Wire), in the 5 L reference	1
Comms	HC-06 Bluetooth, 9600 baud	1
Pumps	NKP-DC-B06B 12 V peristaltic, bidirectional (M1–M4)	4
Motor drivers	L298 2 A modules (2 for pumps, 1 for air)	3
Air pumps	USB 5 V, switched together (ron)	2
Reference reservoir	5 L sealed container (“Wiz tank”)	1
Chambers	200 mL beakers: holding, measurement, 3 M KCl	3
Probe storage	3 M KCl ($\approx 22\%$ w/v)	—

and ref see identical p_{CO_2} at the instant of measurement (the differential’s premise), (ii) blocks room- CO_2 excursions at the source, and (iii) suppresses evaporation. In the final configuration the headspace is ≈ 4 L, the reference ≈ 1 L, and the tank aliquot 100 mL; an air stone placed *in* the reference actively pins the shared p_{CO_2} within minutes (the “anchor”, Sec. 4). Separate or isolated headspaces are explicitly forbidden—two headspaces drift independently and destroy the cancellation.

4 Methods

4.1 Two-phase plateau-following cycle (V4)

Plateau detection is a host-side decision loop, so the PC script drives the firmware step by step. The cycle measures **tank first, then ref**, so the 5 L reference is co-aerated throughout the tank phase and converges quickly when its turn comes:

1. *Prep*: **m3b** (drain KCl from chamber) $\rightarrow$ **m1f** (aquarium $\rightarrow$ holding) $\rightarrow$ **m2f** (holding $\rightarrow$ chamber).
2. *Phase A*: **ron**; measure **tank** every 30 s until plateau (chamber tank + 5 L aerated together).
3. *Transition*: **m2b** parks the tank aliquot back in holding; **m4f** brings reference into the chamber.
4. *Phase B*: **ron**; measure **ref** until plateau (fast—already near equilibrium).
5. **calkh**; then cleanup: **m4b** (recover ref to 5 L), **m1b** (return parked tank to aquarium), **m3f** (restore KCl soak), **airoff** (probe rests in KCl).

4.2 Two-gate plateau criterion

Readings are held as integer milli-pH (mpH) to avoid floating-point boundary jitter. A plateau is declared only when *both* gates hold over a sliding window (Table 2): a short *span* gate (last 4 readings, $\max - \min \leq 2$ pH) rejects noise, and a longer *net* gate (last 8 readings, $|w_{-1} - w_0| \leq 1$ pH) rejects a slow monotonic exponential tail that the span gate alone would mistake for convergence. The longer net look-back (8 vs. 4 points; “B1”) was the fix for early latches in which the tank stopped climbing prematurely. Per-phase wall-clock and read-count caps (**PHASE_MAX** = 7200 s, **MEAS_MAX** = 240) and a consecutive-failure cap (**FAIL_MAX** = 5) bound the loop so it can never hang.

Table 2: Plateau-following parameters (`measure_until_flat`).

Constant	Value	Meaning
FLAT_SPAN_N	4	span window (jitter gate)
FLAT_SPAN_MPH	2	last 4: $\max - \min \leq 2$ mpH
FLAT_NET_N	8	net look-back (drift gate; B1)
FLAT_NET_MPH	1	last 8: $ w_{-1} - w_0 \leq 1$ mpH
MEAS_INTERVAL	30 s	between readings (aeration continuous)
PHASE_MAX_SECS	7200 s	per-phase cap (2 h; 4 h total)
MEAS_MAX	240	per-phase read cap
FAIL_MAX	5	consecutive parse failures tolerated




Figure 2: Measured plateau-following curves (sealed, net-gate run). The tank sample (solid) climbs a slow exponential tail and is latched only when the two-gate criterion (last-4 span ≤ 2 mpH *and* last-8 net ≤ 1 mpH) is met; the co-aerated reference (dashed) is already near equilibrium and latches quickly. The two series sit at different pH because the underlying samples differ (this run predates the null campaign); only their respective *plateau* values enter Eq. (3).

Figure 2 shows a measured pair of approach curves. The tank, freshly drawn into the chamber, climbs a slow first-order tail toward its ceiling and is held by the net gate until it is genuinely flat (latching at ~ 624 s when the 8-point net drift falls to 1 mpH). The reference, having co-aerated in the 5 L reservoir throughout the tank phase, is already at equilibrium and latches in ~ 2.5 min—a direct illustration of why measuring tank-first is faster, and of why the span gate alone (satisfied near ~ 300 s) would have latched the tank early.

4.3 Logging and fault handling

Each completed run appends one line `HH ref_pH tank_pH 8.448 tank_kh temp` to `dkh.dat`. Two distinct markers encode abnormal outcomes: an *all-zeros* line on hard failure (timeout, serial loss, or unverified KCl soak), which additionally *latches*—a subsequent run refuses to measure and rewrites the marker until a human clears it, protecting the probe from unattended pump cycling; and a *negative tank_kh* (magnitude preserved, sign flipped) when a phase hits its cap without reaching plateau, which the data server reads as “did-not-plateau” but which does *not* trip the latch. A *finally* block runs an emergency cleanup (air off, drain chamber, restore KCl soak) on any non-completion.

Table 3: Method/firmware revisions. “ron held” = continuous aeration between runs.

Rev	Period	Change vs. prior
V0	6/14–6/15	First symmetric build; reference drawn from 5 L into a cup; fixed 1200 s aeration.
V1	6/15–6/16	Added continuous aeration between runs (ron held).
V2	6/16–6/17	Reference taken as a small directly-aerated aliquot (holding → chamber); aeration off during the reading.
V3	6/17	Aeration <i>during</i> measurement (pin tank/ref to same $p\text{CO}_2$); path symmetrized; KCl-rinse bubble clean (later dropped).
V4	6/17→	Plateau-following to true equilibrium (no fixed time); two-phase tank-then-ref; anti-hang caps; emergency cleanup. Subsequently: <i>sealed common headspace</i> , <i>5 L active anchor</i> , and the two-gate (B1) criterion.

4.4 Null-test protocol

The diagnostic backbone is the *null test*: the reference and tank vessels are filled with the *same* water so the true alkalinity difference is zero and the correct output is exactly $\Delta\text{pH} = 0$, $\text{dKH} = 8.448$. Two variants are used. In a *partial null*, only the small tank aliquot is replaced with fresh seawater while the aged 5 L reference is retained, so a change in ΔpH isolates the tank water’s contribution. In a *complete null*, reference water is also placed in the tank chamber, so the water difference is zero *by construction* and any residual ΔpH is purely instrumental.

4.5 Revision history and the one-change-at-a-time rule

The build evolved through five revisions (Table 3). Changes were made one causal axis at a time, so that each measurement campaign tests a single mechanism; where two constants were changed together (e.g. stabilization window and convergence band), they were justified as the *same* causal axis (response completeness), not two independent variables.

5 Results

Table 4 lists the full campaign. Throughout, $\Delta\text{pH} < 0$ means the tank read more acidic than the reference (lower apparent alkalinity), and by Eq. (3) the raw error is $\text{dKH} - 8.448 = 8.448(10^{\Delta\text{pH}} - 1)$.

5.1 The dawn- CO_2 error and its diagnosis

The earliest clear pathology was a dawn collapse of ΔpH (run #6, -0.059 ; error -1.09 dKH). The decisive clue was common mode: *both* absolute pH values fell ~ 0.3 overnight (ref $8.109 \rightarrow 7.810$, tank $8.091 \rightarrow 7.751$), reversing the daytime upward trend—the signature of a closed-room CO_2 build-up acidifying the aeration water. A natural experiment confirmed it: on the one ventilated night (run #3, window open) the absolute pH stayed high and ΔpH was a normal $+0.020$, whereas closed nights gave $-0.058/-0.059$. Crucially, room CO_2 is exactly the common-mode variable the differential is *supposed* to cancel; its leakage into ΔpH signalled a broken ref/tank symmetry, not an environmental problem to be ventilated away.

Table 4: Complete measurement campaign (28 runs). “abs pH” is the common-mode level (a room-CO₂ proxy within a single day only). Config abbreviations: S = sealed headspace, A = 5 L active anchor, N = net-gate, B1 = net look-back 8. **Null staging:** a fresh seawater batch is aliquoted at the start; PN = partial null (the 100 mL tank chamber alone is refilled with that fresh aliquot while the aged 5 L reference is kept, isolating the tank-water contribution); CN = complete null (reference water is then transferred into the tank chamber as well, so both vessels hold *identical* water and the water difference is zero by construction).

#	Date/time	Cfg	pH _{ref}	pH _{tank}	Δ pH	dKH	T	Conditions
1	06-14 20:12	V0	7.959	8.021	+0.062	9.741	26.7	first manual symmetric run
2	06-14 21:00	V0	7.997	8.072	+0.075	10.054	26.2	—
3	06-15 05:00	V0	8.057	8.077	+0.020	8.843	26.4	dawn but +; window open (ventilated), low CO ₂
4	06-15 13:00	V1	8.050	8.042	−0.008	8.293	27.0	first continuous aeration
5	06-15 21:00	V1	8.109	8.091	−0.018	8.109	27.3	clean low-CO ₂ point
6	06-16 05:00	V1	7.810	7.751	−0.059	7.362	27.5	dawn closed-room CO ₂ spike (both abs pH fell $\sim$ 0.3)
7	06-16 13:00	V1b	8.091	8.063	−0.028	7.922	27.7	tighter window/band
—	06-16 21:00	—	<i>failed (0.000; power cut, port semaphore)</i>					
8	06-16 21:21	V2	8.042	8.021	−0.021	8.039	28.6	first V2 (manual)
9	06-17 05:00	V2	7.997	7.939	−0.058	7.386	27.8	V2 dawn test <i>failed</i> ; regressed to V1
10	06-17 13:00	V3	8.109	8.091	−0.018	8.095	28.2	first V3; ref under-equilibrated (5 reads)
11	06-17 21:00	V4	8.113	8.071	−0.042	7.669	—	first V4 + trim-mean firmware; clean plateaus
12	06-18 05:00	V4	7.738	7.695	−0.043	7.657	—	decisive: dawn −0.043 $\approx$ evening −0.042 despite 0.4 abs-pH swing $\Rightarrow$ CO ₂ -level independence
13	06-18 13:00	V4	8.081	8.055	−0.026	7.950	28.3	same low CO ₂ as #11 yet −0.026; dKH +0.29 jump $\Rightarrow$ baseline <i>not</i> constant (evaporation, open 200 mL day 5)
14	06-19 05:00	V4+S	7.976	7.968	−0.008	8.283	28.1	seal verif. ①: dawn dip collapses −0.05 $\rightarrow$ −0.008; plateau 33 $\rightarrow$ 4 reads
15	06-19 13:00	V4+S	7.957	7.947	−0.010	8.255	28.6	seal verif. ②: dawn $\rightarrow$ noon jump +0.29 $\rightarrow$ −0.03 ($\sim$ 10 $\times$); evaporation removed
16	06-19 21:00	V4+S	7.969	7.968	−0.001	8.421	26.8	near-perfect null (raw err −0.027, smallest); but +0.166 vs. #14/15 $\Rightarrow$ offset not a perfect constant; AC transient (−1.8°C)
17	06-20 05:00	V4+S	7.922	7.871	−0.051	7.521	27.7	dawn dip recurs; tank 11 $\times$ monotone = early latch (not true plateau); AC cooling slowed kinetics
18	06-20 15:00	V4+SAN	7.975	7.929	−0.046	7.602	27.2	first 5 L anchor; tank <i>true</i> plateau (34 $\times$, 21.6 min) yet −0.046 $\Rightarrow$ real gap , not latch artifact
19	06-20 18:11	V4+S,2 $\times$	7.979	7.924	−0.055	7.443	26.4	controlled exp; 2 $\times$ aeration halves climb (1296 $\rightarrow$ 711 s) but Δ pH unchanged $\Rightarrow$ <i>not</i> kinetics
20	06-20 18:57	V4+S,2 $\times$	7.975	7.931	−0.044	7.636	26.7	ref-first + hose swap; gap persists $\Rightarrow$ order/path/non-equilibrium all rejected
21	06-20 19:50	V4+S	7.975	7.926	−0.049	7.539	26.9	reproduction
22	06-20 20:29	V4+SA	7.977	7.921	−0.056	7.426	26.9	tank-first+park; ref fastest ever (4 $\times$); tank early latch (7.921) $\Rightarrow$ motivates B1
23	06-20 21:00	V4+SA,B1	7.979	7.926	−0.053	7.474	27.0	B1 ok: tank 25 $\times$ /945 s true plateau; confirms #22 was $\sim$ 5 mpH short; gap (7th)
24	06-21 05:00	V4+SA,B1	7.974	7.928	−0.046	7.594	26.7	dawn matches evening cluster $\Rightarrow$ time-of-day dependence ruled out (gap, 8th)
25	06-21 06:00	V4+S,PN	7.973	7.963	−0.010	−8.255	26.7	partial null: fresh tank only; tank pH +37 mpH, Δ pH −0.05 $\rightarrow$ −0.010; tank not plateaued (old cap) $\Rightarrow$ lower bound
26	06-21 13:00	V4+S,PN	7.964	7.956	−0.008	8.298	27.2	partial null, full eq.: dKH 8.298, raw err −0.15; back in good cluster; 4 h cap validated
27	06-21 19:00	V4+S,CN	7.960	7.957	−0.003	8.387	27.4	complete null: reference water into

5.2 Plateau-following and the inter-phase drift conflict

Measuring to a true plateau (V4) removed the CO_2 -level dependence: runs #11/#12 gave $\Delta\text{pH} = -0.042$ and -0.043 across a 0.4-unit swing in absolute pH. But the first V4 run had exposed a competing failure: a 31-minute tank measurement whose pH rose to a peak then *fell* (room CO_2 drifting during the measurement), latching at a value set by whenever the single probe happened to sample each vessel. This *inter-phase drift*, of order

$$\epsilon_{\text{drift}} \approx \Delta t_{\text{tank-ref}} \times \frac{dC}{dt}, \quad (5)$$

sets up a fundamental conflict with a single probe: full equilibrium needs long measurements, but long measurements widen the tank-ref sampling gap Δt and thus the drift. The same form with dT/dt explains run #16, where an air-conditioner cooling overshoot (-1.8°C during the run) coincided with a $+0.166$ dKH excursion.

5.3 Sealing collapses both confounds

A single *sealed common headspace* resolves the conflict at the source. By forcing ref and tank to share one box of air, it makes the common p_{CO_2} intrinsic rather than coincidental, blocks room- CO_2 excursions, and (bonus) suppresses the evaporation that had been concentrating the open tank aliquot. The effect was immediate (runs #14/#15): the dawn dip collapsed from -0.05 to -0.008 ; the dawn $\rightarrow$ noon dKH jump shrank $\sim 10\times$ (from $+0.29$ open to -0.03 sealed); and the plateau time crushed from ~ 33 readings (~ 50 min) to ~ 4 (~ 5 min)—which in turn collapses Δt in Eq. (5), structurally removing the inter-phase drift. The uncorrected error fell from -0.78 to ~ -0.18 dKH.

5.4 The active anchor and the persistence of a real gap

Sealing alone left a slow headspace drift, because a passive shared air volume re-equilibrates only by surface diffusion (tens of minutes). Placing an air stone *in* the 5 L reference (the “anchor”) actively pins the shared p_{CO_2} within minutes; a volume audit confirms the dominance hierarchy (sample $\sim 3\ \mu\text{mol} \ll$ headspace $\sim 65\ \mu\text{mol} \ll$ 5 L $\sim 500\ \mu\text{mol}$). Yet with the anchor and the net-gate in place, a ~ -0.046 to -0.053 gap *persisted* even when the tank reached a verified true plateau (run #18: 34 readings, 21.6 min). A controlled pair of experiments (runs #19/#20) then rejected, in turn, four device-side explanations: doubling the aeration halved the climb time but left ΔpH unchanged (*not* kinetics); reversing the measurement order and swapping hoses did not flip the sign (*not* order, warm-up, or fluid path); and a head-started tank settled to the *same* equilibrium (*not* under-equilibration). The gap was a property of the *water*.

5.5 Null experiments localize the error

Because the device hypotheses were exhausted, we tested the water directly. A *partial null* (run #25, fresh seawater in the tank chamber only) collapsed ΔpH from -0.05 to -0.010 (tank pH $+37$ mpH); at full equilibrium (run #26) it gave dKH = 8.298, raw error -0.15 —back in the good cluster. Swapping fresh water into the tank had recovered almost the entire error, implicating the aged tank aliquot. The *complete null* (run #27), with reference water in *both* vessels so the water difference is zero by construction, returned $\Delta\text{pH} = -0.003$, dKH = 8.387, raw error **-0.061** , with both sides at verified true plateaus despite the 100 mL:5 L volume mismatch—a direct confirmation of volume independence (Sec. 2).

Table 5: Error budget after full diagnosis.

Component	Magnitude	Character / remedy
Aged-tank alkalinity loss	-0.86 dKH	sample handling; use fresh aliquot
Instrument offset	-0.06 dKH	$\Delta\text{pH} = -3$ mpH; below correction value
Precision (1σ)	± 0.08 dKH	electrode noise + residual drift
Room- CO_2 dependence	≈ 0	removed by sealing
Evaporation confound	≈ 0	removed by sealing




Figure 3: Decomposition of the -0.92 dKH raw error. The complete null isolates the instrument offset at -0.06 dKH (one-quarter of the ± 0.25 dKH goal); the partial null attributes the remaining -0.86 dKH to alkalinity loss in the aged tank aliquot (CaCO_3 precipitation / biofilm)—a sample artifact, not an instrument fault.

5.6 Error decomposition and accuracy

The two nulls decompose the headline error (Table 5, Fig. 3):

$$\underbrace{-0.92 \text{ dKH}}_{\text{raw}} = \underbrace{-0.86 \text{ dKH}}_{\substack{\text{aged-tank alkalinity loss} \\ (\text{CaCO}_3 \text{ precip./biofilm})}} + \underbrace{-0.06 \text{ dKH}}_{\substack{\text{instrument offset} \\ (\Delta\text{pH}=-3 \text{ mpH})}}. \quad (6)$$

Across the seven-point cluster (runs #18–24) the *precision* was $\Delta\text{pH} = -0.0499 \pm 0.0048$ (≈ 5 mpH), dKH = 7.531 ± 0.084 —a scatter about one-third of the ± 0.25 dKH goal. With the bulk error identified as a sample artifact and the instrument offset at -0.06 dKH (one-quarter of the goal), the monitor meets ± 0.25 dKH *uncorrected*; zero-point calibration was judged unnecessary, and with normal (non-degraded) reef water the uncorrected error should fall within ± 0.1 dKH.

5.7 Time-resolved breakdown of the null

The decomposition makes a falsifiable prediction: if the bulk error is sample chemistry rather than electronics, an otherwise-perfect null should *decay* as the sample ages. We tested this by re-running the complete null on the same 100 mL aliquot the next day, without touching the instrument (Fig. 4, run #28). The raw error— -0.061 and -0.014 dKH at 0 and 2 h—fell to -0.222 dKH at ~ 14 h ($\Delta\text{pH} = -0.011$; tank pH $7.957 \rightarrow 7.948$ while the larger reference held at 7.959). The drop is well beyond the ± 0.08 dKH precision and was reached at a verified true plateau (tank 59 readings), so it is genuine alkalinity loss, not a latching artifact. The small, exposed aliquot and the unusually long aeration—which raises carbonate saturation and thus drives CaCO_3 precipitation *during* the




Figure 4: Time-resolved breakdown of the complete null. With reference water in *both* vessels (water difference zero by construction), the raw error stays within the ± 0.08 dKH precision (error bars) and the ± 0.25 dKH goal for the first hours, then falls to -0.22 dKH after the small (100 mL), exposed tank aliquot has aged ~ 14 h—a long plateau-following run (59 readings) further driving CaCO_3 precipitation during aeration. The electronics are unchanged; only the sample alkalinity drifts, directly confirming that the dominant error term is sample chemistry (cf. Fig. 3).

measurement—are exactly the worst case for the -0.86 dKH mechanism, and the null broke down accordingly while the instrument offset floor stayed at its fresh value. This is the time-domain confirmation of the error decomposition, and a practical caveat: an anomalously long plateau on a small sample is itself a degradation warning.

6 Discussion

What the diagnosis bought. The instinct on seeing a stable ~ -0.8 dKH offset is to zero-point calibrate. Resisting that—on the methodological principle that a calibration constant should be applied only after independent replication, and never to a small non-independent sample—was decisive, because the offset turned out *not* to be a constant device property (it varied ± 0.085 dKH across the sealed cluster and ultimately proved to be mostly degraded sample water). Calibrating it away would have baked a sample artifact into the instrument and masked the genuine -0.06 dKH floor.

External validation. Our independently derived error structure mirrors published behaviour of the commercial AquaWiz, lending it external support. Community reports identify room- CO_2 variation as the dominant error and a sealed container as the field fix; the manufacturer’s own guidance—that a sharp CO_2 change “only affects readings for about an hour” and that one should “connect the air intake to the skimmer’s outdoor air line”—matches both our ~ 1 -hour inter-phase drift episode and our gas-reference-fixing remedy (we realize it by sealing rather than by an outdoor line). The second-most-reported commercial error, bubble-quality degradation from salt creep (> 2 dKH error, clean every six months), matches our aeration-sufficiency analysis. Reported field accuracy (≤ 0.2 dKH vs. a Hanna checker; ≤ 0.75 dKH uncalibrated over seven weeks; “under 0.5 dKH = excellent, use for trend”) brackets our ± 0.25 dKH result and our trend-oriented intent. Notably, commercial units manage by calibration and cleaning and do not pursue a true equilibrium plateau; our plateau-following step goes one further.

Scientific grounding. The kinetic separation underlying the method is well established: the aqueous CO_2 hydration/dehydration relaxation completes in ~ 16 s [2], so chemistry is never rate-limiting; the slow step is physical gas exchange, ~ 15 min for direct bubbling [5], consistent with our observed plateau times. The differential “ A_T from p_{CO_2} -equilibration plus pH” principle has a recent analytical-chemistry basis demonstrated down to 0.5 mL samples at $\pm 1\text{--}2\ \mu\text{mol kg}^{-1}$ precision [6], supporting small-volume miniaturization as the correct direction.

Limitations. (i) With a single probe the inter-phase drift of Eq. (5) cannot be driven to zero; it is only suppressed (by sealing and by minimizing Δt). (ii) The dominant historical error was sample degradation, so the operating discipline—use a fresh, adequately sized tank aliquot and recirculate rather than store—is part of the instrument, not an afterthought. (iii) Long plateaus under slow drift can make a measurement run to hours; the wall-clock caps bound this but a hard external scheduler limit must be set above the code’s 4-hour ceiling. (iv) Absolute pH is informative only as a within-day CO_2 proxy; only ΔpH is quantitative.

Recommended operation. Seal the common headspace; anchor the reference with active aeration; measure tank-then-ref with the two-gate plateau criterion; use fresh sample water; clean the air path periodically; and treat the output as a trend instrument calibrated, if at all, against an independent kit on a monthly cadence keyed to the ref $\leftrightarrow$ tank divergence.

7 Conclusion

A reagent-free, single-probe differential-pH aeration monitor can measure reef carbonate alkalinity to better than ± 0.25 dKH *without* zero-point calibration, once three design rules are observed: a single sealed common headspace, plateau-following to true gas equilibrium, and an actively aerated reference anchor. The value of the eight-day campaign was less in any one number than in *decomposing* a -0.92 dKH raw error: null experiments showed that -0.86 dKH of it was alkalinity loss in a degraded sample aliquot and only -0.06 dKH was the instrument. The differential design delivered its promised common-mode rejection of room CO_2 and probe drift; what remained was chemistry in a beaker, not error in the electronics. We expect the design rules and the null-decomposition methodology to transfer to other low-cost alkalinity instruments.

Reproducibility

Firmware, the host measurement script (`measure_kh_once.py`), the raw log (`dkh.dat`), and the full measurement ledger are available in the project repository [1].

Acknowledgments

This project was completed thanks to the passion of Mango-miyeok of REEF FAMILY and all the participants of the DIY meetup.

References

- [1] nowforyou (REEF FAMILY), “ReefKeeper: open differential-pH aeration alkalinity monitor (firmware, host scripts, measurement ledger),” project repository, <https://github.com/taeseokyi/reefkeeper>.
- [2] K. S. Johnson, “Carbon dioxide hydration and dehydration kinetics in seawater,” *Limnol. Oceanogr.* **27**(5), 849–855 (1982). doi:10.4319/lo.1982.27.5.0849.
- [3] A. L. Soli and R. H. Byrne, “CO₂ system hydration and dehydration kinetics and the equilibrium CO₂/H₂CO₃ ratio in aqueous NaCl solution,” *Mar. Chem.* (2002).
- [4] R. E. Zeebe and D. Wolf-Gladrow, *CO₂ in Seawater: Equilibrium, Kinetics, Isotopes*, Elsevier Oceanography Series 65 (2001).
- [5] T. Takahashi, “Methods for measurement of $p\text{CO}_2$ and total CO₂ in seawater” (carbonate-system methodology; direct-bubbling equilibration ~ 15 min).
- [6] K. Fleger, X. Liu, W. Berelson, J. Adkins *et al.*, “Total alkalinity measurements in small samples based on CO₂ equilibration and spectrophotometric pH,” *Anal. Chim. Acta* (2025). doi:10.1016/j.aca.2025.344432.
- [7] (Gas-stripping first-order equilibration; $\ln(C_0/C_t) \propto \phi/V$.) Roberts, *Geophys. Res. Lett.* (2005).
- [8] C. S. Egleston, C. L. Sabine, F. M. M. Morel, “Revelle (buffer) factor of the seawater carbonate system,” *Global Biogeochem. Cycles* (2010).
- [9] Reef-hobbyist community field reports on the commercial AquaWiz alkalinity monitor (room-CO₂ error, sealed-container fix, bubble degradation, manufacturer guidance): Reef2Reef threads #1159272, #1073924, #1115439; UltimateReef #939234; Bay Area Reefers #40433. (Sources recorded as forum thread identifiers; quotations are translations of community/manufacture posts and should be verified against the originals before formal publication.)